

Only weight-level unlearning rides the probe axis

Llama-3-8B · $\cos(\text{refusal-removal direction, probe } w)$ at last prompt token

negative = activations move toward harmless / comply ($-w$)

